

# **LLM Ecosystems: From Foundations to Frontiers**

ACSPRI Summer Program 2026 Dr. Maria Prokofieva



# Why Do LLMs Matter in Social Science?

---

## LLMs scale human-level language understanding

They can analyze millions of words — reports, interviews, articles — with nuanced pattern recognition that mimics expert judgment at scale.

## They reduce cost + time for text-heavy workflows

Manual coding, thematic tagging, or policy trend analysis that once took weeks can now be done in hours using reproducible, scriptable methods.

## Applications include policy, legal analysis, qualitative coding

LLMs are already used to monitor citizen sentiment, identify emerging legal themes, cluster qualitative survey data, and flag ethical risks in reports.

## Academic access is now feasible — no longer just for Big Tech

Open-weight models and APIs (like Hugging Face, OpenAI, Gemini) now allow researchers to apply cutting-edge models with minimal infrastructure.

## What is an LLM?

---

Machine learning models  
trained on billions of words

Work via token prediction,  
not memorization

Understand and generate  
human-like text

Used via APIs or code

Can summarize, classify,  
extract, and generate

Not perfect, but powerful  
and improving

## LLMs Power Text Analysis

---

### NLP


*Natural Language  
Processing*

Auto-code  
interviews or  
survey responses

Summarize long  
documents for  
policymaking

Detect sentiment  
shifts across large  
datasets

Classify thousands of  
news articles in  
minutes



## How Do LLMs Actually Work?

---



### Trained on billions of words

Ingests books, articles, websites — huge corpus



### Uses transformer architecture

“Attention” lets model focus on relevant input



### Predicts tokens, not

### retrieves facts

Generates next word based on probability



### Understands context in chunks (context window)

But not global memory — limited to recent input

## Principles

---

### Self-Attention

Key idea: understand word relevance across a sequence

### Positional Encoding

Inject structure into unordered tokens

### Transformer Block

Composed of dense layers + normalization + residual paths, and stacked deep

# Major Types of LLM Models

## Encoder-only

**Goal:** Understand input text deeply

**What it does:** Reads the full input and builds a contextual representation

**Popular examples:** BERT, RoBERTa, DistilBERT

**Used for:**

Text classification

Named Entity Recognition

Sentence similarity *No text generation*

## Decoder-only

**Goal:** Generate fluent text from left to right

**What it does:** Predicts one token at a time based on previous tokens

**Popular examples:** GPT-2, GPT-3, GPT-4, Claude, Gemini

**Used for:**

Chatbots

Text generation

Summarization *These models don't "see" the entire input at once — they build it as they go.*

## Encoder–Decoder (Seq2Seq)

**Goal:** Transform one sequence into another

**What it does:** Encoder reads the full input

Decoder generates output


**Popular examples:** T5, FLAN-T5, BART, mT5

**Used for:**

Translation

Text summarization

Question answering *They're the most flexible, but more complex.*



## Major Types of LLM Models

# Modalities in LLMs

---

## Modality

*What kinds of input  
can LLMs understand?*



### Text

Natural Language  
input/output



### Image

Visual input: images/photos



### Audio

Voice/sound input/output



### Video

Moving sequences



### Multilingual

Cross-lingual understanding




### Structured Data

Tables, JSON, code

## How Modalities Map to Model Types

---



# Foundational vs Proprietary Models

## Foundational Models


- Trained on broad, general-purpose datasets (e.g. Common Crawl, Wikipedia, books)
- Released as base models.
- Examples: LLaMA 3, Mistral, Falcon.



## Proprietary Models

- Developed by private companies
- High performance, but API-only access
- Customization often limited.
- Examples: GPT-4, Gemini, Claude.





## Model Families Landscape

# Model= Architecture + Weights

*What Makes a Model Work?  
Architecture gives **structure**.  
Weights give **skill**.*

## Architecture

**Blueprint:** How many layers, what type, how they connect

Defines **how**:

- data flows through the model
- data patterns can be learned

*Architecture choice determines if a model can handle images (CNNs) or text (transformers).*

## Weights

**Learned values:** What the model knows after training

**Numerical values** (e.g., 0.347, -1.892) that scale input data.

*Weights determine if a model actually works well for your specific task.*



## Proprietary Models: What You Don't Get to See

---



*Understanding the hidden layers behind closed-source LLMs*

- **Weights are not released** — only accessible via API.
- **Training data** is undisclosed .
- **Model architecture** may not disclosed.
- **Fine-tuning access** is limited or expensive.
- **Usage tracking + safety layers** often run in the background.

# Parameters

*What does a trillion-parameter model actually mean?*

## Numerical values

Parameters are the ~~that define how a neural network processes~~ input data to produce output.

## Not "knowledge"


That's the weights and patterns formed by parameters during training.

## Not "rules"

The model has no explicit logic—just math.

## Not fixed

They change during training via optimization.



# What Parameters Actually Do

## Transform Input Data

Every input (text, image, etc.) is converted to numbers (vectors). Parameters are the coefficients in the functions that: **Scale (weight)**, **Shift (bias)**, or **Gate (attention)** these numbers.



## Enable Learning

During training, parameters are adjusted via:  
**Loss calculation:** How wrong the model's output is.  
**Gradient descent:** Tweaking parameters to reduce loss.



## Dictate Model Behavior

The arrangement and values of parameters determine:  
**Capacity:** Can it memorize complex patterns?  
**Generalization:** Can it apply patterns to new data?



💡 *Parameters = the knobs, levers, and memory of a model. They shape how input becomes output — and how models evolve.*

## **What Does “Open Weights” Actually Mean?**

*Open Weights ≠ Full Openness*

### **Downloadable model weights**

Training data (usually unreleased or partially described)

### **Ability to run/fine-tune locally**

Full license to deploy commercially (varies by model)

### **Architecture details (often published)**

Tokenizers, eval metrics, or guardrails often separate

# Does Bigger = Better?

*More parameters ≠ better results — unless trained well and aligned to the right task.*

## Memorization vs Generalization



### Bigger ≠ Smarter

Large models can **memorize training data** if not trained properly — making them brittle or biased.

**Example:** A poorly trained 70B model might repeat toxic forum posts instead of understanding nuance in a new context.

## Finetuned > Gigantic




### Smaller + Targeted = Better

A smaller model, **finetuned on clinical trials**, can outperform GPT-4 on biomedical summarization.

**Example:** BioGPT or a custom Gemini finetune often beats generic LLMs in specialized settings.

## Cost Reality




### More Params = More Compute = More \$\$\$

Every billion parameters adds cost — for both training and inference.

**Example:** GPT-4-turbo is cheaper than GPT-4 because it optimizes how those parameters are used — not because it's smaller.

## Preview — From Parameters to Meaning (Day 2)



*"The model doesn't just generate text. It learns to position language in space."*